

predict_YOLOv3

May 1, 2020

0.1 This code is adapted from [How to implement a YOLO \(v3\) object detector from scratch in PyTorch](#) and its implementation, [Github repo](#).

1 1. Import

1.1 1) General

```
[1]: import time
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import Variable
import numpy as np
import cv2
import argparse
import os
import os.path as osp
import pickle as pkl
import pandas as pd
import random
```

1.2 2) Model

```
[2]: class Darknet(nn.Module):
    def __init__(self, cfgfile):
        super(Darknet, self).__init__()
        self.blocks = parse_cfg(cfgfile)
        self.net_info, self.module_list = create_modules(self.blocks)

    def forward(self, x, CUDA):
        modules = self.blocks[1:]
        outputs = {} #We cache the outputs for the route layer

        write = 0
        for i, module in enumerate(modules):
            module_type = (module["type"])

            if module_type == "convolutional" or module_type == "upsample":
```

```

        x = self.module_list[i](x)

    elif module_type == "route":
        layers = module["layers"]
        layers = [int(a) for a in layers]

        if (layers[0]) > 0:
            layers[0] = layers[0] - i

        if len(layers) == 1:
            x = outputs[i + (layers[0])]

        else:
            if (layers[1]) > 0:
                layers[1] = layers[1] - i

            map1 = outputs[i + layers[0]]
            map2 = outputs[i + layers[1]]
            x = torch.cat((map1, map2), 1)

    elif module_type == "shortcut":
        from_ = int(module["from"])
        x = outputs[i-1] + outputs[i+from_]

    elif module_type == 'yolo':
        anchors = self.module_list[i][0].anchors
        #Get the input dimensions
        inp_dim = int (self.net_info["height"])

        #Get the number of classes
        num_classes = int (module["classes"])

        #Transform
        x = x.data
        x = predict_transform(x, inp_dim, anchors, num_classes, CUDA)
        if not write: #if no collector has been intialised.

            detections = x
            write = 1

        else:
            detections = torch.cat((detections, x), 1)

    outputs[i] = x

return detections

```

```

def load_weights(self, weightfile):
    #Open the weights file
    fp = open(weightfile, "rb")

    #The first 5 values are header information
    # 1. Major version number
    # 2. Minor Version Number
    # 3. Subversion number
    # 4,5. Images seen by the network (during training)
    header = np.fromfile(fp, dtype = np.int32, count = 5)
    self.header = torch.from_numpy(header)
    self.seen = self.header[3]

    weights = np.fromfile(fp, dtype = np.float32)

    ptr = 0
    for i in range(len(self.module_list)):
        module_type = self.blocks[i + 1]["type"]

        #If module_type is convolutional load weights
        #Otherwise ignore.

        if module_type == "convolutional":
            model = self.module_list[i]
            try:
                batch_normalize = int(self.blocks[i+1]["batch_normalize"])
            except:
                batch_normalize = 0

            conv = model[0]

            if (batch_normalize):
                bn = model[1]

                #Get the number of weights of Batch Norm Layer
                num_bn_biases = bn.bias.numel()

                #Load the weights
                bn_biases = torch.from_numpy(weights[ptr:ptr +
↪num_bn_biases])
                ptr += num_bn_biases

                bn_weights = torch.from_numpy(weights[ptr: ptr +
↪num_bn_biases])
                ptr += num_bn_biases

```

```

        bn_running_mean = torch.from_numpy(weights[ptr: ptr + ↵
↵num_bn_biases])
        ptr += num_bn_biases

        bn_running_var = torch.from_numpy(weights[ptr: ptr + ↵
↵num_bn_biases])
        ptr += num_bn_biases

        #Cast the loaded weights into dims of model weights.
        bn_biases = bn_biases.view_as(bn.bias.data)
        bn_weights = bn_weights.view_as(bn.weight.data)
        bn_running_mean = bn_running_mean.view_as(bn.running_mean)
        bn_running_var = bn_running_var.view_as(bn.running_var)

        #Copy the data to model
        bn.bias.data.copy_(bn_biases)
        bn.weight.data.copy_(bn_weights)
        bn.running_mean.copy_(bn_running_mean)
        bn.running_var.copy_(bn_running_var)

    else:
        #Number of biases
        num_biases = conv.bias.numel()

        #Load the weights
        conv_biases = torch.from_numpy(weights[ptr: ptr + ↵
↵num_biases])

        ptr = ptr + num_biases

        #reshape the loaded weights according to the dims of the ↵
↵model weights
        conv_biases = conv_biases.view_as(conv.bias.data)

        #Finally copy the data
        conv.bias.data.copy_(conv_biases)

        #Let us load the weights for the Convolutional layers
        num_weights = conv.weight.numel()

        #Do the same as above for weights
        conv_weights = torch.from_numpy(weights[ptr:ptr+num_weights])
        ptr = ptr + num_weights

        conv_weights = conv_weights.view_as(conv.weight.data)
        conv.weight.data.copy_(conv_weights)

```

```

[3]: def predict_transform(prediction, inp_dim, anchors, num_classes, CUDA = True):

    batch_size = prediction.size(0)
    stride = inp_dim // prediction.size(2)
    grid_size = inp_dim // stride
    bbox_attrs = 5 + num_classes
    num_anchors = len(anchors)

    prediction = prediction.view(batch_size, bbox_attrs*num_anchors,
    ↪grid_size*grid_size)
    prediction = prediction.transpose(1,2).contiguous()
    prediction = prediction.view(batch_size, grid_size*grid_size*num_anchors,
    ↪bbox_attrs)
    anchors = [(a[0]/stride, a[1]/stride) for a in anchors]

    #Sigmoid the centre_X, centre_Y. and object confidence
    prediction[:, :, 0] = torch.sigmoid(prediction[:, :, 0])
    prediction[:, :, 1] = torch.sigmoid(prediction[:, :, 1])
    prediction[:, :, 4] = torch.sigmoid(prediction[:, :, 4])

    #Add the center offsets
    grid = np.arange(grid_size)
    a,b = np.meshgrid(grid, grid)

    x_offset = torch.FloatTensor(a).view(-1,1)
    y_offset = torch.FloatTensor(b).view(-1,1)

    if CUDA:
        x_offset = x_offset.cuda()
        y_offset = y_offset.cuda()

    x_y_offset = torch.cat((x_offset, y_offset), 1).repeat(1,num_anchors).
    ↪view(-1,2).unsqueeze(0)

    prediction[:, :, :2] += x_y_offset

    #log space transform height and the width
    anchors = torch.FloatTensor(anchors)

    if CUDA:
        anchors = anchors.cuda()

    anchors = anchors.repeat(grid_size*grid_size, 1).unsqueeze(0)
    prediction[:, :, 2:4] = torch.exp(prediction[:, :, 2:4])*anchors

    prediction[:, :, 5: 5 + num_classes] = torch.sigmoid((prediction[:, :, 5 : 5 +
    ↪num_classes]))

```

```

prediction[:, :, :4] *= stride

return prediction

```

```

[4]: class EmptyLayer(nn.Module):
      def __init__(self):
          super(EmptyLayer, self).__init__()

      class DetectionLayer(nn.Module):
          def __init__(self, anchors):
              super(DetectionLayer, self).__init__()
              self.anchors = anchors

```

```

[5]: def parse_cfg(cfgfile):
      """
      Takes a configuration file

      Returns a list of blocks. Each blocks describes a block in the neural
      network to be built. Block is represented as a dictionary in the list

      """
      file = open(cfgfile, 'r')
      lines = file.read().split('\n')           # store the lines in
      ↪ a list
      lines = [x for x in lines if len(x) > 0]   # get read of the
      ↪ empty lines
      lines = [x for x in lines if x[0] != '#']   # get rid of comments
      lines = [x.rstrip().lstrip() for x in lines] # get rid of fringe
      ↪ whitespaces

      block = {}
      blocks = []

      for line in lines:
          if line[0] == "[":                     # This marks the start of a new block
              if len(block) != 0:                 # If block is not empty, implies it is
              ↪ storing values of previous block.
                  blocks.append(block)           # add it the blocks list
                  block = {}                     # re-init the block
                  block["type"] = line[1:-1].rstrip()
              else:
                  key,value = line.split("=")
                  block[key.rstrip()] = value.lstrip()
              blocks.append(block)

      return blocks

```

```

def create_modules(blocks):
    net_info = blocks[0]      #Captures the information about the input and
    →pre-processing
    module_list = nn.ModuleList()
    prev_filters = 3
    output_filters = []

    for index, x in enumerate(blocks[1:]):
        module = nn.Sequential()

        #check the type of block
        #create a new module for the block
        #append to module_list

        if (x["type"] == "convolutional"):
            #Get the info about the layer
            activation = x["activation"]
            try:
                batch_normalize = int(x["batch_normalize"])
                bias = False
            except:
                batch_normalize = 0
                bias = True

            filters= int(x["filters"])
            padding = int(x["pad"])
            kernel_size = int(x["size"])
            stride = int(x["stride"])

            if padding:
                pad = (kernel_size - 1) // 2
            else:
                pad = 0

            #Add the convolutional layer
            conv = nn.Conv2d(prev_filters, filters, kernel_size, stride, pad,
    →bias = bias)
            module.add_module("conv_{0}".format(index), conv)

            #Add the Batch Norm Layer
            if batch_normalize:
                bn = nn.BatchNorm2d(filters)
                module.add_module("batch_norm_{0}".format(index), bn)

            #Check the activation.
            #It is either Linear or a Leaky ReLU for YOLO

```

```

        if activation == "leaky":
            activn = nn.LeakyReLU(0.1, inplace = True)
            module.add_module("leaky_{0}".format(index), activn)

        #If it's an upsampling layer
        #We use Bilinear2dUpsampling
        elif (x["type"] == "upsample"):
            stride = int(x["stride"])
            upsample = nn.Upsample(scale_factor = 2, mode = "bilinear")
            module.add_module("upsample_{0}".format(index), upsample)

        #If it is a route layer
        elif (x["type"] == "route"):
            x["layers"] = x["layers"].split(',')
            #Start of a route
            start = int(x["layers"][0])
            #end, if there exists one.
            try:
                end = int(x["layers"][1])
            except:
                end = 0
            #Positive anotation
            if start > 0:
                start = start - index
            if end > 0:
                end = end - index
            route = EmptyLayer()
            module.add_module("route_{0}".format(index), route)
            if end < 0:
                filters = output_filters[index + start] + output_filters[index_
→+ end]

        else:
            filters= output_filters[index + start]

        #shortcut corresponds to skip connection
        elif x["type"] == "shortcut":
            shortcut = EmptyLayer()
            module.add_module("shortcut_{0}".format(index), shortcut)

        elif x["type"] == "yolo":
            mask = x["mask"].split(",")
            mask = [int(x) for x in mask]

            anchors = x["anchors"].split(",")
            anchors = [int(a) for a in anchors]
            anchors = [(anchors[i], anchors[i+1]) for i in range(0,
→len(anchors),2)]
            anchors = [anchors[i] for i in mask]

```



```

        detection = DetectionLayer(anchors)
        module.add_module("Detection_{}".format(index), detection)
        module_list.append(module)
        prev_filters = filters
        output_filters.append(filters)
    return (net_info, module_list)

```

1.3 3) etc

```

[6]: def bbox_iou(box1, box2):
    """
    Returns the IoU of two bounding boxes

    """
    #Get the coordinates of bounding boxes
    b1_x1, b1_y1, b1_x2, b1_y2 = box1[:,0], box1[:,1], box1[:,2], box1[:,3]
    b2_x1, b2_y1, b2_x2, b2_y2 = box2[:,0], box2[:,1], box2[:,2], box2[:,3]

    #get the corrdinates of the intersection rectangle
    inter_rect_x1 = torch.max(b1_x1, b2_x1)
    inter_rect_y1 = torch.max(b1_y1, b2_y1)
    inter_rect_x2 = torch.min(b1_x2, b2_x2)
    inter_rect_y2 = torch.min(b1_y2, b2_y2)

    #Intersection area
    inter_area = torch.clamp(inter_rect_x2 - inter_rect_x1 + 1, min=0) * torch.
    ↪ clamp(inter_rect_y2 - inter_rect_y1 + 1, min=0)

    #Union Area
    b1_area = (b1_x2 - b1_x1 + 1)*(b1_y2 - b1_y1 + 1)
    b2_area = (b2_x2 - b2_x1 + 1)*(b2_y2 - b2_y1 + 1)

    iou = inter_area / (b1_area + b2_area - inter_area)

    return iou

```

```

[7]: def write_results(prediction, confidence, num_classes, nms_conf = 0.4):
    conf_mask = (prediction[:, :, 4] > confidence).float().unsqueeze(2)
    prediction = prediction*conf_mask

    box_corner = prediction.new(prediction.shape)
    box_corner[:, :, 0] = (prediction[:, :, 0] - prediction[:, :, 2])/2
    box_corner[:, :, 1] = (prediction[:, :, 1] - prediction[:, :, 3])/2
    box_corner[:, :, 2] = (prediction[:, :, 0] + prediction[:, :, 2])/2
    box_corner[:, :, 3] = (prediction[:, :, 1] + prediction[:, :, 3])/2

```

```

prediction[:, :, :4] = box_corner[:, :, :4]

batch_size = prediction.size(0)

write = False

for ind in range(batch_size):
    image_pred = prediction[ind]          #image Tensor
    #confidence thresholding
    #NMS

    max_conf, max_conf_score = torch.max(image_pred[:, 5:5+ num_classes], 1)
    max_conf = max_conf.float().unsqueeze(1)
    max_conf_score = max_conf_score.float().unsqueeze(1)
    seq = (image_pred[:, :5], max_conf, max_conf_score)
    image_pred = torch.cat(seq, 1)

    non_zero_ind = (torch.nonzero(image_pred[:, 4]))
    try:
        image_pred_ = image_pred[non_zero_ind.squeeze(), :].view(-1, 7)
    except:
        continue

    if image_pred_.shape[0] == 0:
        continue
#

    #Get the various classes detected in the image
    img_classes = unique(image_pred_[:, -1]) # -1 index holds the class
    → index

    for cls in img_classes:
        #perform NMS
        #get the detections with one particular class
        cls_mask = image_pred_*(image_pred_[:, -1] == cls).float().
        →unsqueeze(1)
        class_mask_ind = torch.nonzero(cls_mask[:, -2]).squeeze()
        image_pred_class = image_pred_[class_mask_ind].view(-1, 7)

        #sort the detections such that the entry with the maximum objectness
        #confidence is at the top
        conf_sort_index = torch.sort(image_pred_class[:, 4], descending =
        →True ) [1]
        image_pred_class = image_pred_class[conf_sort_index]

```

```

        idx = image_pred_class.size(0)    #Number of detections

        for i in range(idx):
            #Get the IOUs of all boxes that come after the one we are
            → looking at
            #in the loop
            try:
                ious = bbox_iou(image_pred_class[i].unsqueeze(0),
            → image_pred_class[i+1:])
                except ValueError:
                    break

                except IndexError:
                    break

            #Zero out all the detections that have IoU > treshhold
            iou_mask = (ious < nms_conf).float().unsqueeze(1)
            image_pred_class[i+1:] *= iou_mask

            #Remove the non-zero entries
            non_zero_ind = torch.nonzero(image_pred_class[:,4]).squeeze()
            image_pred_class = image_pred_class[non_zero_ind].view(-1,7)

        batch_ind = image_pred_class.new(image_pred_class.size(0), 1).
        → fill_(ind)    #Repeat the batch_id for as many detections of the class cls
        → in the image
        seq = batch_ind, image_pred_class

        if not write:
            output = torch.cat(seq,1)
            write = True
        else:
            out = torch.cat(seq,1)
            output = torch.cat((output,out))

        try:
            return output
        except:
            return 0

```

```

[8]: def load_classes(namesfile):
    fp = open(namesfile, "r")
    names = fp.read().split("\n")[:-1]
    return names

```

2 2. Run

```
[9]: img_dir = 'testset-img'
det_dir = 'predicted_boxes'
batch_size = 1
confidence = 0.5
nms_thesh = 0.4
cfgfile = 'cfg/yolov3.cfg'
weightsfile = 'weights/yolov3.weights'
start = 0
reso = 416
CUDA = torch.cuda.is_available()
num_classes = 80      #For COCO
classes = load_classes("data/coco.names")
```

```
[10]: #Set up the neural network
print("Loading network.....")
model = Darknet(cfgfile)
model.load_weights(weightsfile)
print("Network successfully loaded")

model.net_info["height"] = reso
inp_dim = int(model.net_info["height"])
assert inp_dim % 32 == 0
assert inp_dim > 32

#If there's a GPU available, put the model on GPU
if CUDA:
    model.cuda()

#Set the model in evaluation mode
model.eval()
```

Loading network...

Network successfully loaded

```
[10]: Darknet(
  (module_list): ModuleList(
    (0): Sequential(
      (conv_0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (batch_norm_0): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_0): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (1): Sequential(
      (conv_1): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1,
```

```

1), bias=False)
    (batch_norm_1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leaky_1): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (2): Sequential(
      (conv_2): Conv2d(64, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_2): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (3): Sequential(
      (conv_3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (batch_norm_3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_3): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (4): Sequential(
      (shortcut_4): EmptyLayer()
    )
    (5): Sequential(
      (conv_5): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
      (batch_norm_5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_5): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (6): Sequential(
      (conv_6): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_6): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_6): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (7): Sequential(
      (conv_7): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (batch_norm_7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_7): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (8): Sequential(
      (shortcut_8): EmptyLayer()
    )
    (9): Sequential(
      (conv_9): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_9): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    (leaky_9): LeakyReLU(negative_slope=0.1, inplace=True)
)
(10): Sequential(
  (conv_10): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_10): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_10): LeakyReLU(negative_slope=0.1, inplace=True)
)
(11): Sequential(
  (shortcut_11): EmptyLayer()
)
(12): Sequential(
  (conv_12): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
  (batch_norm_12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_12): LeakyReLU(negative_slope=0.1, inplace=True)
)
(13): Sequential(
  (conv_13): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_13): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_13): LeakyReLU(negative_slope=0.1, inplace=True)
)
(14): Sequential(
  (conv_14): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_14): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_14): LeakyReLU(negative_slope=0.1, inplace=True)
)
(15): Sequential(
  (shortcut_15): EmptyLayer()
)
(16): Sequential(
  (conv_16): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_16): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_16): LeakyReLU(negative_slope=0.1, inplace=True)
)
(17): Sequential(
  (conv_17): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_17): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (leaky_17): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (18): Sequential(
      (shortcut_18): EmptyLayer()
    )
    (19): Sequential(
      (conv_19): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_19): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_19): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (20): Sequential(
      (conv_20): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (batch_norm_20): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_20): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (21): Sequential(
      (shortcut_21): EmptyLayer()
    )
    (22): Sequential(
      (conv_22): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_22): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_22): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (23): Sequential(
      (conv_23): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (batch_norm_23): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_23): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (24): Sequential(
      (shortcut_24): EmptyLayer()
    )
    (25): Sequential(
      (conv_25): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_25): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_25): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (26): Sequential(
      (conv_26): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
      (batch_norm_26): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    (leaky_26): LeakyReLU(negative_slope=0.1, inplace=True)
)
(27): Sequential(
  (shortcut_27): EmptyLayer()
)
(28): Sequential(
  (conv_28): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_28): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_28): LeakyReLU(negative_slope=0.1, inplace=True)
)
(29): Sequential(
  (conv_29): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_29): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_29): LeakyReLU(negative_slope=0.1, inplace=True)
)
(30): Sequential(
  (shortcut_30): EmptyLayer()
)
(31): Sequential(
  (conv_31): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_31): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_31): LeakyReLU(negative_slope=0.1, inplace=True)
)
(32): Sequential(
  (conv_32): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_32): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_32): LeakyReLU(negative_slope=0.1, inplace=True)
)
(33): Sequential(
  (shortcut_33): EmptyLayer()
)
(34): Sequential(
  (conv_34): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_34): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_34): LeakyReLU(negative_slope=0.1, inplace=True)
)
(35): Sequential(
  (conv_35): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)

```



```

        (batch_norm_35): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_35): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (36): Sequential(
        (shortcut_36): EmptyLayer()
    )
    (37): Sequential(
        (conv_37): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (batch_norm_37): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_37): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (38): Sequential(
        (conv_38): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_38): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_38): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (39): Sequential(
        (conv_39): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_39): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_39): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (40): Sequential(
        (shortcut_40): EmptyLayer()
    )
    (41): Sequential(
        (conv_41): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_41): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_41): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (42): Sequential(
        (conv_42): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_42): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_42): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (43): Sequential(
        (shortcut_43): EmptyLayer()
    )
    (44): Sequential(

```

```

        (conv_44): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_44): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_44): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (45): Sequential(
        (conv_45): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_45): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_45): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (46): Sequential(
        (shortcut_46): EmptyLayer()
    )
    (47): Sequential(
        (conv_47): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_47): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_47): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (48): Sequential(
        (conv_48): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_48): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_48): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (49): Sequential(
        (shortcut_49): EmptyLayer()
    )
    (50): Sequential(
        (conv_50): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_50): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_50): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (51): Sequential(
        (conv_51): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_51): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_51): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (52): Sequential(
        (shortcut_52): EmptyLayer()
    )

```

```

(53): Sequential(
  (conv_53): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_53): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_53): LeakyReLU(negative_slope=0.1, inplace=True)
)
(54): Sequential(
  (conv_54): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_54): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_54): LeakyReLU(negative_slope=0.1, inplace=True)
)
(55): Sequential(
  (shortcut_55): EmptyLayer()
)
(56): Sequential(
  (conv_56): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_56): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_56): LeakyReLU(negative_slope=0.1, inplace=True)
)
(57): Sequential(
  (conv_57): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_57): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_57): LeakyReLU(negative_slope=0.1, inplace=True)
)
(58): Sequential(
  (shortcut_58): EmptyLayer()
)
(59): Sequential(
  (conv_59): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_59): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_59): LeakyReLU(negative_slope=0.1, inplace=True)
)
(60): Sequential(
  (conv_60): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batch_norm_60): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_60): LeakyReLU(negative_slope=0.1, inplace=True)
)
(61): Sequential(
  (shortcut_61): EmptyLayer()
)

```

```

    )
    (62): Sequential(
      (conv_62): Conv2d(512, 1024, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
      (batch_norm_62): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_62): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (63): Sequential(
      (conv_63): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_63): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_63): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (64): Sequential(
      (conv_64): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_64): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_64): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (65): Sequential(
      (shortcut_65): EmptyLayer()
    )
    (66): Sequential(
      (conv_66): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_66): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_66): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (67): Sequential(
      (conv_67): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_67): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_67): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (68): Sequential(
      (shortcut_68): EmptyLayer()
    )
    (69): Sequential(
      (conv_69): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_69): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

        (leaky_69): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (70): Sequential(
      (conv_70): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_70): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_70): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (71): Sequential(
      (shortcut_71): EmptyLayer()
    )
    (72): Sequential(
      (conv_72): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_72): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_72): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (73): Sequential(
      (conv_73): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_73): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_73): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (74): Sequential(
      (shortcut_74): EmptyLayer()
    )
    (75): Sequential(
      (conv_75): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_75): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_75): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (76): Sequential(
      (conv_76): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_76): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leaky_76): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (77): Sequential(
      (conv_77): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_77): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    (leaky_77): LeakyReLU(negative_slope=0.1, inplace=True)
)
(78): Sequential(
  (conv_78): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_78): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_78): LeakyReLU(negative_slope=0.1, inplace=True)
)
(79): Sequential(
  (conv_79): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_79): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_79): LeakyReLU(negative_slope=0.1, inplace=True)
)
(80): Sequential(
  (conv_80): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_80): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_80): LeakyReLU(negative_slope=0.1, inplace=True)
)
(81): Sequential(
  (conv_81): Conv2d(1024, 255, kernel_size=(1, 1), stride=(1, 1))
)
(82): Sequential(
  (Detection_82): DetectionLayer()
)
(83): Sequential(
  (route_83): EmptyLayer()
)
(84): Sequential(
  (conv_84): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_84): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leaky_84): LeakyReLU(negative_slope=0.1, inplace=True)
)
(85): Sequential(
  (upsample_85): Upsample(scale_factor=2.0, mode=bilinear)
)
(86): Sequential(
  (route_86): EmptyLayer()
)
(87): Sequential(
  (conv_87): Conv2d(768, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

```

```

        (batch_norm_87): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_87): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (88): Sequential(
        (conv_88): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_88): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_88): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (89): Sequential(
        (conv_89): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_89): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_89): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (90): Sequential(
        (conv_90): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_90): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_90): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (91): Sequential(
        (conv_91): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_91): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_91): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (92): Sequential(
        (conv_92): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batch_norm_92): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_92): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (93): Sequential(
        (conv_93): Conv2d(512, 255, kernel_size=(1, 1), stride=(1, 1))
    )
    (94): Sequential(
        (Detection_94): DetectionLayer()
    )
    (95): Sequential(
        (route_95): EmptyLayer()
    )
    (96): Sequential(

```

```

        (conv_96): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_96): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_96): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (97): Sequential(
        (upsample_97): Upsample(scale_factor=2.0, mode=bilinear)
    )
    (98): Sequential(
        (route_98): EmptyLayer()
    )
    (99): Sequential(
        (conv_99): Conv2d(384, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batch_norm_99): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_99): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (100): Sequential(
        (conv_100): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (batch_norm_100): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_100): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (101): Sequential(
        (conv_101): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (batch_norm_101): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_101): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (102): Sequential(
        (conv_102): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
        (batch_norm_102): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_102): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (103): Sequential(
        (conv_103): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
        (batch_norm_103): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leaky_103): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (104): Sequential(
        (conv_104): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),

```



```
padding=(1, 1), bias=False)
    (batch_norm_104): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leaky_104): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (105): Sequential(
      (conv_105): Conv2d(256, 255, kernel_size=(1, 1), stride=(1, 1))
    )
    (106): Sequential(
      (Detection_106): DetectionLayer()
    )
  )
)
```

```
[11]: from os import listdir
from os.path import isfile, join

filename = [f for f in listdir(img_dir)]
imlist = [join(img_dir, f) for f in listdir(img_dir) if isfile(join(img_dir,
→f)) and f.endswith('.jpg')]
loaded_ims = [cv2.imread(x) for x in imlist]
```

```
[12]: def letterbox_image(img, inp_dim):
    '''resize image with unchanged aspect ratio using padding'''
    img_w, img_h = img.shape[1], img.shape[0]
    w, h = inp_dim
    new_w = int(img_w * min(w/img_w, h/img_h))
    new_h = int(img_h * min(w/img_w, h/img_h))
    resized_image = cv2.resize(img, (new_w,new_h), interpolation = cv2.
→INTER_CUBIC)

    canvas = np.full((inp_dim[1], inp_dim[0], 3), 128)

    canvas[(h-new_h)//2:(h-new_h)//2 + new_h,(w-new_w)//2:(w-new_w)//2 + new_w,
→:] = resized_image

    return canvas

def prep_image(img, inp_dim):
    '''
    Prepare image for inputting to the neural network.

    Returns a Variable
    '''
    img = (letterbox_image(img, (inp_dim, inp_dim)))
    img = img[:, :, :-1].transpose((2,0,1)).copy()
    img = torch.from_numpy(img).float().div(255.0).unsqueeze(0)
```

```

    return img

def unique(tensor):
    tensor_np = tensor.cpu().numpy()
    unique_np = np.unique(tensor_np)
    unique_tensor = torch.from_numpy(unique_np)

    tensor_res = tensor.new(unique_tensor.shape)
    tensor_res.copy_(unique_tensor)
    return tensor_res

```

```

[13]: im_batches = list(map(prepare_image, loaded_imgs, [inp_dim for x in
    ↪range(len(imlist))]))

```

```

#List containing dimensions of original images
im_dim_list = [(x.shape[1], x.shape[0]) for x in loaded_imgs]
im_dim_list = torch.FloatTensor(im_dim_list).repeat(1,2)

```

```

[14]: p = model(Variable(im_batches[0], volatile = True), CUDA)

```

<ipython-input-14-17ae41444f36>:1: UserWarning: volatile was removed and now has no effect. Use `with torch.no_grad():` instead.

```

p = model(Variable(im_batches[0], volatile = True), CUDA)
/Users/chanho/miniconda3/envs/eval/lib/python3.8/site-
packages/torch/nn/functional.py:2970: UserWarning: Default upsampling behavior
when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify
align_corners=True if the old behavior is desired. See the documentation of
nn.Upsample for details.

```

```

    warnings.warn("Default upsampling behavior when mode={} is changed ")

```

```

[15]: write = 0
start_det_loop = time.time()
for i, batch in enumerate(im_batches):
    #     print(i)
    start = time.time()
    prediction = model(Variable(batch, volatile = True), CUDA)
    prediction = write_results(prediction, confidence, num_classes, nms_conf =
    ↪nms_thesh)
    end = time.time()
    #     print(prediction)

    if type(prediction) == int:
        for im_num, image in enumerate(imlist[i*batch_size: min((i +
    ↪1)*batch_size, len(imlist))]):
            im_id = i*batch_size + im_num
            print("{0:20s} predicted in {1:6.3f} seconds".format(image.split("/")
    ↪)[-1], (end - start)/batch_size))

```

```

        print("{0:20s} {1:s}".format("Objects Detected:", ""))
        print("-----")
        continue

    prediction[:,0] += i*batch_size    #transform the attribute from index in
    ↪ batch to index in imlist

    if not write:                      #If we have't initialised output
        output = prediction
        write = 1
    else:
        output = torch.cat((output,prediction))
    for im_num, image in enumerate(imlist[i*batch_size: min((i + 1
    ↪ 1)*batch_size, len(imlist))]):
        im_id = i*batch_size + im_num
        objs = [classes[int(x[-1])] for x in output if int(x[0]) == im_id]
        print("{0:20s} predicted in {1:6.3f} seconds".format(image.split("/")
    ↪)[-1], (end - start)/batch_size))
        print("{0:20s} {1:s}".format("Objects Detected:", " ".join(objs)))
        print("-----")

```

<ipython-input-15-19b1c4e5e187>:6: UserWarning: volatile was removed and now has no effect. Use `with torch.no_grad():` instead.

```

prediction = model(Variable(batch, volatile = True), CUDA)
/Users/distiller/project/conda/conda-
bld/pytorch_1587428077867/work/torch/csrc/utils/python_arg_parser.cpp:756:

```

UserWarning: This overload of nonzero is deprecated:

```

    nonzero(Tensor input, *, Tensor out)

```

Consider using one of the following signatures instead:

```

    nonzero(Tensor input, *, bool as_tuple)

```

```

frame0080.jpg      predicted in  0.732 seconds

```

```

Objects Detected:  person person
-----

```

```

frame0094.jpg      predicted in  0.754 seconds

```

```

Objects Detected:  person person person person person truck
-----

```

```

frame0043.jpg      predicted in  0.724 seconds

```

```

Objects Detected:  person person person person person person person person
person
-----

```

```

frame0057.jpg      predicted in  0.757 seconds

```

```

Objects Detected:  person person car truck
-----

```

```

frame0056.jpg      predicted in  0.734 seconds

```

```

Objects Detected:  car car
-----

```

```

frame0042.jpg      predicted in  0.719 seconds

```

Objects Detected:	person car car

frame0095.jpg	predicted in 0.717 seconds
Objects Detected:	person person person person

frame0081.jpg	predicted in 0.722 seconds
Objects Detected:	person

frame0097.jpg	predicted in 0.723 seconds
Objects Detected:	person person person person person

frame0083.jpg	predicted in 0.713 seconds
Objects Detected:	person

frame0068.jpg	predicted in 0.715 seconds
Objects Detected:	

frame0054.jpg	predicted in 0.713 seconds
Objects Detected:	person person car car

frame0040.jpg	predicted in 0.714 seconds
Objects Detected:	person person person person person person person car

frame0041.jpg	predicted in 0.745 seconds
Objects Detected:	person person person person person person person person person
person	

frame0055.jpg	predicted in 0.719 seconds
Objects Detected:	person person person

frame0069.jpg	predicted in 0.703 seconds
Objects Detected:	

frame0082.jpg	predicted in 0.701 seconds
Objects Detected:	car truck

frame0096.jpg	predicted in 0.710 seconds
Objects Detected:	person person person person person person person backpack

frame0092.jpg	predicted in 0.704 seconds
Objects Detected:	person person person

frame0086.jpg	predicted in 0.712 seconds
Objects Detected:	person person person car

frame0051.jpg	predicted in 0.715 seconds
Objects Detected:	person person person handbag

frame0045.jpg	predicted in 0.713 seconds
Objects Detected:	person person person person

frame0079.jpg	predicted in 0.701 seconds
Objects Detected:	person person

frame0078.jpg	predicted in 0.709 seconds
Objects Detected:	

frame0044.jpg	predicted in 0.706 seconds
Objects Detected:	person person person person person person person

frame0050.jpg	predicted in 0.707 seconds
Objects Detected:	person person person person person person person person handbag

frame0087.jpg	predicted in 0.700 seconds
Objects Detected:	person person person person

frame0093.jpg	predicted in 0.703 seconds
Objects Detected:	

frame0085.jpg	predicted in 0.703 seconds
Objects Detected:	

frame0091.jpg	predicted in 0.702 seconds
Objects Detected:	

frame0046.jpg	predicted in 0.704 seconds
Objects Detected:	person person person person person person person

frame0052.jpg	predicted in 0.712 seconds
Objects Detected:	person person person person person person

frame0053.jpg	predicted in 0.730 seconds
Objects Detected:	person person person person person person person person person person

frame0047.jpg	predicted in 0.735 seconds
Objects Detected:	person person person person person person person bus

frame0090.jpg	predicted in 0.758 seconds
Objects Detected:	person person person person car car

frame0084.jpg	predicted in 0.721 seconds
Objects Detected:	

frame0008.jpg	predicted in 0.719 seconds
Objects Detected:	person person person person person person person person person

person person person

frame0020.jpg predicted in 0.717 seconds

Objects Detected: car car

frame0034.jpg predicted in 0.712 seconds

Objects Detected: person person person person person person person

frame0035.jpg predicted in 0.711 seconds

Objects Detected: person person person person

frame0021.jpg predicted in 0.708 seconds

Objects Detected:

frame0009.jpg predicted in 0.704 seconds

Objects Detected:

frame0037.jpg predicted in 0.704 seconds

Objects Detected: person person person person person person person

frame0023.jpg predicted in 0.714 seconds

Objects Detected: person person person person person person person car car
truck

frame0022.jpg predicted in 0.713 seconds

Objects Detected: person person person person person person person person

frame0036.jpg predicted in 0.707 seconds

Objects Detected: car

frame0032.jpg predicted in 0.730 seconds

Objects Detected: person person person person person person person handbag

frame0026.jpg predicted in 0.734 seconds

Objects Detected: person person person person handbag

frame0027.jpg predicted in 0.706 seconds

Objects Detected: person car

frame0033.jpg predicted in 0.712 seconds

Objects Detected: person person person

frame0025.jpg predicted in 0.727 seconds

Objects Detected: person person person person person person person person
person person person person person

frame0031.jpg predicted in 0.717 seconds

Objects Detected:

```

-----
frame0019.jpg      predicted in  0.767 seconds
Objects Detected:  person person person person person person person person
-----
frame0018.jpg      predicted in  0.763 seconds
Objects Detected:  person person person person person person person car truck
-----
frame0030.jpg      predicted in  0.723 seconds
Objects Detected:  car car
-----
frame0024.jpg      predicted in  0.713 seconds
Objects Detected:  car
-----
frame0029.jpg      predicted in  0.715 seconds
Objects Detected:  person person car truck
-----
frame0001.jpg      predicted in  0.716 seconds
Objects Detected:  person person person person person person person person
-----
frame0015.jpg      predicted in  0.719 seconds
Objects Detected:  person person person person person person person person
-----
frame0014.jpg      predicted in  0.710 seconds
Objects Detected:  person car car
-----
frame0000.jpg      predicted in  0.709 seconds
Objects Detected:  person person person person person person person bicycle
-----
frame0028.jpg      predicted in  0.724 seconds
Objects Detected:  person person person person person person person person car
-----
frame0016.jpg      predicted in  0.716 seconds
Objects Detected:  person person person
-----
frame0002.jpg      predicted in  0.728 seconds
Objects Detected:  person person person person person person person person person
-----
frame0003.jpg      predicted in  0.727 seconds
Objects Detected:  person person person person person person person person person
person person person person
-----
frame0017.jpg      predicted in  0.730 seconds
Objects Detected:  person person car car
-----
frame0013.jpg      predicted in  0.718 seconds
Objects Detected:  person person person person person person car
-----
frame0007.jpg      predicted in  0.720 seconds

```

Objects Detected:	person person person person person person person

frame0006.jpg	predicted in 0.706 seconds
Objects Detected:	person car

frame0012.jpg	predicted in 0.707 seconds
Objects Detected:	person car

frame0004.jpg	predicted in 0.705 seconds
Objects Detected:	person person person person person person person person

frame0010.jpg	predicted in 0.714 seconds
Objects Detected:	person person person backpack

frame0038.jpg	predicted in 0.715 seconds
Objects Detected:	person person person

frame0039.jpg	predicted in 0.713 seconds
Objects Detected:	person person person person person person

frame0011.jpg	predicted in 0.705 seconds
Objects Detected:	person person person

frame0005.jpg	predicted in 0.724 seconds
Objects Detected:	person person

frame0089.jpg	predicted in 0.732 seconds
Objects Detected:	person person car truck

frame0062.jpg	predicted in 0.746 seconds
Objects Detected:	person person person person person person person person

frame0076.jpg	predicted in 0.726 seconds
Objects Detected:	person person person person person car

frame0077.jpg	predicted in 0.731 seconds
Objects Detected:	person person person person person person

frame0063.jpg	predicted in 0.741 seconds
Objects Detected:	person person person person person

frame0088.jpg	predicted in 0.725 seconds
Objects Detected:	person person person person person bus truck

frame0049.jpg	predicted in 0.735 seconds
Objects Detected:	person person person person person person person car

frame0075.jpg	predicted in 0.737 seconds

Objects Detected:

```
-----
frame0061.jpg      predicted in  0.733 seconds
Objects Detected:  person person person person person person person
-----
frame0060.jpg      predicted in  0.734 seconds
Objects Detected:  person person person person person person
-----
frame0074.jpg      predicted in  0.732 seconds
Objects Detected:  person person person person person
-----
frame0048.jpg      predicted in  0.723 seconds
Objects Detected:  person person person
-----
frame0070.jpg      predicted in  0.733 seconds
Objects Detected:  person person person
-----
frame0064.jpg      predicted in  0.747 seconds
Objects Detected:  person person person
-----
frame0058.jpg      predicted in  0.755 seconds
Objects Detected:  person person
-----
frame0059.jpg      predicted in  0.762 seconds
Objects Detected:  person person person person person person
-----
frame0065.jpg      predicted in  0.830 seconds
Objects Detected:  person person person person person person person
person person
-----
frame0071.jpg      predicted in  0.750 seconds
Objects Detected:  car car truck
-----
frame0098.jpg      predicted in  0.719 seconds
Objects Detected:  person person person person person person person
-----
frame0067.jpg      predicted in  0.706 seconds
Objects Detected:  person person person person person person person
-----
frame0073.jpg      predicted in  0.707 seconds
Objects Detected:  person person person
-----
frame0072.jpg      predicted in  0.721 seconds
Objects Detected:  person person
-----
frame0066.jpg      predicted in  0.714 seconds
Objects Detected:  person truck
-----
```

```
frame0099.jpg          predicted in 0.726 seconds
Objects Detected:      person person person
-----
```

```
[16]: im_dim_list = torch.index_select(im_dim_list, 0, output[:,0].long())

scaling_factor = torch.min(inp_dim/im_dim_list,1)[0].view(-1,1)

output[:,[1,3]] -= (inp_dim - scaling_factor*im_dim_list[:,0].view(-1,1))/2
output[:,[2,4]] -= (inp_dim - scaling_factor*im_dim_list[:,1].view(-1,1))/2

output[:,1:5] /= scaling_factor
```

```
[17]: def write(x):
        img_id = int(x[0])
        xmin, ymin, xmax, ymax = int(x[1]), int(x[2]), int(x[3]), int(x[4])
        conf_score = float(x[6])
        cls = int(x[7])
        label = "{0}".format(classes[cls])
        return filename[img_id], xmin, ymin, xmax, ymax, conf_score, label
```

```
[18]: df = pd.DataFrame(list(map(write, output)), columns=['fn', 'xmin', 'ymin', 'xmax', 'ymax', 'conf_score', 'label'])
df.head()
```

```
[18]:
```

	fn	xmin	ymin	xmax	ymax	conf_score	label
0	frame0080.jpg	512	345	534	398	0.993586	person
1	frame0080.jpg	755	320	789	393	0.987500	person
2	frame0094.jpg	812	294	852	399	0.999899	person
3	frame0094.jpg	955	312	1005	394	0.999720	person
4	frame0094.jpg	533	235	594	392	0.999767	person

```
[19]: import copy

class_dict = {}
for l in classes:
    class_dict[l] = {}
    for f in filename:
        df_l_f = df[(df.label == l) & (df.fn == f)].values.tolist()
        if df_l_f:
            class_dict[l][f] = {}
            class_dict[l][f]['boxes'] = []
            class_dict[l][f]['scores'] = []

            for d in df_l_f:
                class_dict[l][f]['boxes'].append([d[1], d[2], d[3], d[4]])
                class_dict[l][f]['scores'].append(d[5])
```

```
[22]: import json

for l in ['person', 'car']:
    with open(det_dir+'/predicted_boxes-YOLOv3-'+l+'.json', 'w') as fp:
        json.dump(class_dict[l], fp)
```