

# How to implement a YOLO (v3) object detector from scratch in PyTorch

## [Part 1: Understanding How YOLO works \(https://blog.paperspace.com/how-to-implement-a-yolo-object-detector-in-pytorch/\)](https://blog.paperspace.com/how-to-implement-a-yolo-object-detector-in-pytorch/)

Github repo ([https://github.com/ayooshkathuria/YOLO\\_v3\\_tutorial\\_from\\_scratch](https://github.com/ayooshkathuria/YOLO_v3_tutorial_from_scratch))

## [Part 2 : Creating the layers of the network architecture \(https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-2/\)](https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-2/)

```
In [1]: # mkdir cfg
        # cd cfg
        # wget https://raw.githubusercontent.com/pjreddie/darknet/master/cfg/yolov3.cfg
```

```
In [2]: from __future__ import division

import time
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import Variable
import numpy as np
import cv2
import argparse
import os
import os.path as osp
import pickle as pkl
import pandas as pd
import random
```

```
In [3]: def parse_cfg(cfgfile):
        """
        Takes a configuration file

        Returns a list of blocks. Each blocks describes a block in the
        neural
        network to be built. Block is represented as a dictionary in th
        e list

        """
        file = open(cfgfile, 'r')
        lines = file.read().split('\n')           # store
        the lines in a list
        lines = [x for x in lines if len(x) > 0]   # get re
        ad of the empty lines
        lines = [x for x in lines if x[0] != '#']   # get ri
        d of comments
        lines = [x.rstrip().lstrip() for x in lines] # get ri
        d of fringe whitespaces

        block = {}
        blocks = []

        for line in lines:
            if line[0] == "[":                     # This marks the start of
            a new block
                if len(block) != 0:                 # If block is not empty, i
            mples it is storing values of previous block.
                    blocks.append(block)           # add it the blocks list
                    block = {}                     # re-init the block
                    block["type"] = line[1:-1].rstrip()
            else:
                key,value = line.split("=")
                block[key.rstrip()] = value.lstrip()
            blocks.append(block)

        return blocks
```

```
In [4]: class EmptyLayer(nn.Module):
        def __init__(self):
            super(EmptyLayer, self).__init__()
```

```
In [5]: class DetectionLayer(nn.Module):
        def __init__(self, anchors):
            super(DetectionLayer, self).__init__()
            self.anchors = anchors
```

```
In [6]: def create_modules(blocks):
        net_info = blocks[0]           #Captures the information about the in
        put and pre-processing
        module_list = nn.ModuleList()
```

```

prev_filters = 3
output_filters = []

for index, x in enumerate(blocks[1:]):
    module = nn.Sequential()

    #check the type of block
    #create a new module for the block
    #append to module_list

    if (x["type"] == "convolutional"):
        #Get the info about the layer
        activation = x["activation"]
        try:
            batch_normalize = int(x["batch_normalize"])
            bias = False
        except:
            batch_normalize = 0
            bias = True

        filters= int(x["filters"])
        padding = int(x["pad"])
        kernel_size = int(x["size"])
        stride = int(x["stride"])

        if padding:
            pad = (kernel_size - 1) // 2
        else:
            pad = 0

        #Add the convolutional layer
        conv = nn.Conv2d(prev_filters, filters, kernel_size, st
ride, pad, bias = bias)
        module.add_module("conv_{0}".format(index), conv)

        #Add the Batch Norm Layer
        if batch_normalize:
            bn = nn.BatchNorm2d(filters)
            module.add_module("batch_norm_{0}".format(index), b
n)

        #Check the activation.
        #It is either Linear or a Leaky ReLU for YOLO
        if activation == "leaky":
            activn = nn.LeakyReLU(0.1, inplace = True)
            module.add_module("leaky_{0}".format(index), activn
)

        #If it's an upsampling layer
        #We use Bilinear2dUpsampling
        elif (x["type"] == "upsample"):
            stride = int(x["stride"])
            upsample = nn.Upsample(scale_factor = 2, mode = "biline

```

```

ar")
        module.add_module("upsample_{}".format(index), upsample
    )
    #If it is a route layer
    elif (x["type"] == "route"):
        x["layers"] = x["layers"].split(',')
        #Start of a route
        start = int(x["layers"][0])
        #end, if there exists one.
        try:
            end = int(x["layers"][1])
        except:
            end = 0
        #Positive anotation
        if start > 0:
            start = start - index
        if end > 0:
            end = end - index
        route = EmptyLayer()
        module.add_module("route_{0}".format(index), route)
        if end < 0:
            filters = output_filters[index + start] + output_filters[index + end]
        else:
            filters= output_filters[index + start]

    #shortcut corresponds to skip connection
    elif x["type"] == "shortcut":
        shortcut = EmptyLayer()
        module.add_module("shortcut_{}".format(index), shortcut
    )

    elif x["type"] == "yolo":
        mask = x["mask"].split(",")
        mask = [int(x) for x in mask]

        anchors = x["anchors"].split(",")
        anchors = [int(a) for a in anchors]
        anchors = [(anchors[i], anchors[i+1]) for i in range(0, len(anchors), 2)]
        anchors = [anchors[i] for i in mask]

        detection = DetectionLayer(anchors)
        module.add_module("Detection_{}".format(index), detection)

    module_list.append(module)
    prev_filters = filters
    output_filters.append(filters)
    return (net_info, module_list)

```

```

In [7]: blocks = parse_cfg("cfg/yolov3.cfg")
        print(create_modules(blocks))

```

```

({'type': 'net', 'batch': '64', 'subdivisions': '16', 'width': '41

```

```

6', 'height': '416', 'channels': '3', 'momentum': '0.9', 'decay':
'0.0005', 'angle': '0', 'saturation': '1.5', 'exposure': '1.5', 'hue': '.1', 'learning_rate': '0.001', 'burn_in': '1000', 'max_batches': '500200', 'policy': 'steps', 'steps': '400000,450000', 'scales': '.1,.1'}, ModuleList(
  (0): Sequential(
    (conv_0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (batch_norm_0): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_0): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (1): Sequential(
    (conv_1): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (batch_norm_1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_1): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (2): Sequential(
    (conv_2): Conv2d(64, 32, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batch_norm_2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_2): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (3): Sequential(
    (conv_3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (batch_norm_3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_3): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (4): Sequential(
    (shortcut_4): EmptyLayer()
  )
  (5): Sequential(
    (conv_5): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (batch_norm_5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_5): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (6): Sequential(
    (conv_6): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batch_norm_6): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_6): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (7): Sequential(
    (conv_7): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

```

```
(batch_norm_7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_7): LeakyReLU(negative_slope=0.1, inplace=True)
)
(8): Sequential(
  (shortcut_8): EmptyLayer()
)
(9): Sequential(
  (conv_9): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_9): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_9): LeakyReLU(negative_slope=0.1, inplace=True)
)
(10): Sequential(
  (conv_10): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_10): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_10): LeakyReLU(negative_slope=0.1, inplace=True)
)
(11): Sequential(
  (shortcut_11): EmptyLayer()
)
(12): Sequential(
  (conv_12): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
  (batch_norm_12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_12): LeakyReLU(negative_slope=0.1, inplace=True)
)
(13): Sequential(
  (conv_13): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_13): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_13): LeakyReLU(negative_slope=0.1, inplace=True)
)
(14): Sequential(
  (conv_14): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_14): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_14): LeakyReLU(negative_slope=0.1, inplace=True)
)
(15): Sequential(
  (shortcut_15): EmptyLayer()
)
(16): Sequential(
  (conv_16): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_16): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
```

```
        (leaky_16): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (17): Sequential(
      (conv_17): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_17): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_17): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (18): Sequential(
      (shortcut_18): EmptyLayer()
    )
    (19): Sequential(
      (conv_19): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_19): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_19): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (20): Sequential(
      (conv_20): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_20): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_20): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (21): Sequential(
      (shortcut_21): EmptyLayer()
    )
    (22): Sequential(
      (conv_22): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_22): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_22): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (23): Sequential(
      (conv_23): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_23): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_23): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (24): Sequential(
      (shortcut_24): EmptyLayer()
    )
    (25): Sequential(
      (conv_25): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_25): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_25): LeakyReLU(negative_slope=0.1, inplace=True)
    )
  )
)
```

```
(26): Sequential(
  (conv_26): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_26): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_26): LeakyReLU(negative_slope=0.1, inplace=True)
)
(27): Sequential(
  (shortcut_27): EmptyLayer()
)
(28): Sequential(
  (conv_28): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_28): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_28): LeakyReLU(negative_slope=0.1, inplace=True)
)
(29): Sequential(
  (conv_29): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_29): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_29): LeakyReLU(negative_slope=0.1, inplace=True)
)
(30): Sequential(
  (shortcut_30): EmptyLayer()
)
(31): Sequential(
  (conv_31): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_31): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_31): LeakyReLU(negative_slope=0.1, inplace=True)
)
(32): Sequential(
  (conv_32): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_32): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_32): LeakyReLU(negative_slope=0.1, inplace=True)
)
(33): Sequential(
  (shortcut_33): EmptyLayer()
)
(34): Sequential(
  (conv_34): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_34): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_34): LeakyReLU(negative_slope=0.1, inplace=True)
)
(35): Sequential(
  (conv_35): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
```



```
padding=(1, 1), bias=False)
    (batch_norm_35): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_35): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (36): Sequential(
    (shortcut_36): EmptyLayer()
  )
  (37): Sequential(
    (conv_37): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (batch_norm_37): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_37): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (38): Sequential(
    (conv_38): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batch_norm_38): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_38): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (39): Sequential(
    (conv_39): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (batch_norm_39): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_39): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (40): Sequential(
    (shortcut_40): EmptyLayer()
  )
  (41): Sequential(
    (conv_41): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batch_norm_41): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_41): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (42): Sequential(
    (conv_42): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (batch_norm_42): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_42): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (43): Sequential(
    (shortcut_43): EmptyLayer()
  )
  (44): Sequential(
    (conv_44): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batch_norm_44): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine
```

```
ine=True, track_running_stats=True)
    (leaky_44): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (45): Sequential(
    (conv_45): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (batch_norm_45): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
    (leaky_45): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (46): Sequential(
    (shortcut_46): EmptyLayer()
  )
  (47): Sequential(
    (conv_47): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (batch_norm_47): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
    (leaky_47): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (48): Sequential(
    (conv_48): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (batch_norm_48): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
    (leaky_48): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (49): Sequential(
    (shortcut_49): EmptyLayer()
  )
  (50): Sequential(
    (conv_50): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (batch_norm_50): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
    (leaky_50): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (51): Sequential(
    (conv_51): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
    (batch_norm_51): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
    (leaky_51): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (52): Sequential(
    (shortcut_52): EmptyLayer()
  )
  (53): Sequential(
    (conv_53): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
    (batch_norm_53): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
    (leaky_53): LeakyReLU(negative_slope=0.1, inplace=True)
```

```
)
(54): Sequential(
  (conv_54): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_54): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_54): LeakyReLU(negative_slope=0.1, inplace=True)
)
(55): Sequential(
  (shortcut_55): EmptyLayer()
)
(56): Sequential(
  (conv_56): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_56): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_56): LeakyReLU(negative_slope=0.1, inplace=True)
)
(57): Sequential(
  (conv_57): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_57): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_57): LeakyReLU(negative_slope=0.1, inplace=True)
)
(58): Sequential(
  (shortcut_58): EmptyLayer()
)
(59): Sequential(
  (conv_59): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_59): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_59): LeakyReLU(negative_slope=0.1, inplace=True)
)
(60): Sequential(
  (conv_60): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_60): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_60): LeakyReLU(negative_slope=0.1, inplace=True)
)
(61): Sequential(
  (shortcut_61): EmptyLayer()
)
(62): Sequential(
  (conv_62): Conv2d(512, 1024, kernel_size=(3, 3), stride=(2, 2)
, padding=(1, 1), bias=False)
  (batch_norm_62): BatchNorm2d(1024, eps=1e-05, momentum=0.1, af
fine=True, track_running_stats=True)
  (leaky_62): LeakyReLU(negative_slope=0.1, inplace=True)
)
(63): Sequential(
```

```
(conv_63): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1)
, bias=False)
(batch_norm_63): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(leaky_63): LeakyReLU(negative_slope=0.1, inplace=True)
)
(64): Sequential(
  (conv_64): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), bias=False)
  (batch_norm_64): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_64): LeakyReLU(negative_slope=0.1, inplace=True)
)
(65): Sequential(
  (shortcut_65): EmptyLayer()
)
(66): Sequential(
  (conv_66): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1)
, bias=False)
  (batch_norm_66): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_66): LeakyReLU(negative_slope=0.1, inplace=True)
)
(67): Sequential(
  (conv_67): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), bias=False)
  (batch_norm_67): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_67): LeakyReLU(negative_slope=0.1, inplace=True)
)
(68): Sequential(
  (shortcut_68): EmptyLayer()
)
(69): Sequential(
  (conv_69): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1)
, bias=False)
  (batch_norm_69): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_69): LeakyReLU(negative_slope=0.1, inplace=True)
)
(70): Sequential(
  (conv_70): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), bias=False)
  (batch_norm_70): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_70): LeakyReLU(negative_slope=0.1, inplace=True)
)
(71): Sequential(
  (shortcut_71): EmptyLayer()
)
(72): Sequential(
  (conv_72): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1)
, bias=False)
```

```
(batch_norm_72): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_72): LeakyReLU(negative_slope=0.1, inplace=True)
)
(73): Sequential(
  (conv_73): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_73): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_73): LeakyReLU(negative_slope=0.1, inplace=True)
)
(74): Sequential(
  (shortcut_74): EmptyLayer()
)
(75): Sequential(
  (conv_75): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_75): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_75): LeakyReLU(negative_slope=0.1, inplace=True)
)
(76): Sequential(
  (conv_76): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_76): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_76): LeakyReLU(negative_slope=0.1, inplace=True)
)
(77): Sequential(
  (conv_77): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_77): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_77): LeakyReLU(negative_slope=0.1, inplace=True)
)
(78): Sequential(
  (conv_78): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_78): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_78): LeakyReLU(negative_slope=0.1, inplace=True)
)
(79): Sequential(
  (conv_79): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_79): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_79): LeakyReLU(negative_slope=0.1, inplace=True)
)
(80): Sequential(
  (conv_80): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_80): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_80): LeakyReLU(negative_slope=0.1, inplace=True)
)
```

```

fine=True, track_running_stats=True)
    (leaky_80): LeakyReLU(negative_slope=0.1, inplace=True)
)
(81): Sequential(
  (conv_81): Conv2d(1024, 255, kernel_size=(1, 1), stride=(1, 1)
)
)
(82): Sequential(
  (Detection_82): DetectionLayer()
)
(83): Sequential(
  (route_83): EmptyLayer()
)
(84): Sequential(
  (conv_84): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_84): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_84): LeakyReLU(negative_slope=0.1, inplace=True)
)
(85): Sequential(
  (upsample_85): Upsample(scale_factor=2.0, mode=bilinear)
)
(86): Sequential(
  (route_86): EmptyLayer()
)
(87): Sequential(
  (conv_87): Conv2d(768, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_87): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_87): LeakyReLU(negative_slope=0.1, inplace=True)
)
(88): Sequential(
  (conv_88): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_88): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_88): LeakyReLU(negative_slope=0.1, inplace=True)
)
(89): Sequential(
  (conv_89): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
  (batch_norm_89): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_89): LeakyReLU(negative_slope=0.1, inplace=True)
)
(90): Sequential(
  (conv_90): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
  (batch_norm_90): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
  (leaky_90): LeakyReLU(negative_slope=0.1, inplace=True)
)

```

```

    )
    (91): Sequential(
      (conv_91): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_91): BatchNorm2d(256, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
      (leaky_91): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (92): Sequential(
      (conv_92): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_92): BatchNorm2d(512, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
      (leaky_92): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (93): Sequential(
      (conv_93): Conv2d(512, 255, kernel_size=(1, 1), stride=(1, 1))
    )
    (94): Sequential(
      (Detection_94): DetectionLayer()
    )
    (95): Sequential(
      (route_95): EmptyLayer()
    )
    (96): Sequential(
      (conv_96): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_96): BatchNorm2d(128, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
      (leaky_96): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (97): Sequential(
      (upsample_97): Upsample(scale_factor=2.0, mode=bilinear)
    )
    (98): Sequential(
      (route_98): EmptyLayer()
    )
    (99): Sequential(
      (conv_99): Conv2d(384, 128, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_99): BatchNorm2d(128, eps=1e-05, momentum=0.1, aff
ine=True, track_running_stats=True)
      (leaky_99): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (100): Sequential(
      (conv_100): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), bias=False)
      (batch_norm_100): BatchNorm2d(256, eps=1e-05, momentum=0.1, af
fine=True, track_running_stats=True)
      (leaky_100): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (101): Sequential(
      (conv_101): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1)

```

```

, bias=False)
    (batch_norm_101): BatchNorm2d(128, eps=1e-05, momentum=0.1, af
fine=True, track_running_stats=True)
    (leaky_101): LeakyReLU(negative_slope=0.1, inplace=True)
)
    (102): Sequential(
      (conv_102): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), bias=False)
      (batch_norm_102): BatchNorm2d(256, eps=1e-05, momentum=0.1, af
fine=True, track_running_stats=True)
      (leaky_102): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (103): Sequential(
      (conv_103): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1)
, bias=False)
      (batch_norm_103): BatchNorm2d(128, eps=1e-05, momentum=0.1, af
fine=True, track_running_stats=True)
      (leaky_103): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (104): Sequential(
      (conv_104): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1)
, padding=(1, 1), bias=False)
      (batch_norm_104): BatchNorm2d(256, eps=1e-05, momentum=0.1, af
fine=True, track_running_stats=True)
      (leaky_104): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (105): Sequential(
      (conv_105): Conv2d(256, 255, kernel_size=(1, 1), stride=(1, 1)
)
    )
    (106): Sequential(
      (Detection_106): DetectionLayer()
    )
  ))

```

### [Part 3 : Implementing the the forward pass of the network \(https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-3/\)](https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-3/)

```

In [8]: class Darknet(nn.Module):
        def __init__(self, cfgfile):
            super(Darknet, self).__init__()
            self.blocks = parse_cfg(cfgfile)
            self.net_info, self.module_list = create_modules(self.block
s)

        def forward(self, x, CUDA):
            modules = self.blocks[1:]
            outputs = {} #We cache the outputs for the route layer

```



```

write = 0
for i, module in enumerate(modules):
    module_type = (module["type"])

    if module_type == "convolutional" or module_type == "upsample":
        x = self.module_list[i](x)

    elif module_type == "route":
        layers = module["layers"]
        layers = [int(a) for a in layers]

        if (layers[0]) > 0:
            layers[0] = layers[0] - i

        if len(layers) == 1:
            x = outputs[i + (layers[0])]

        else:
            if (layers[1]) > 0:
                layers[1] = layers[1] - i

            map1 = outputs[i + layers[0]]
            map2 = outputs[i + layers[1]]
            x = torch.cat((map1, map2), 1)

    elif module_type == "shortcut":
        from_ = int(module["from"])
        x = outputs[i-1] + outputs[i+from_]

    elif module_type == 'yolo':
        anchors = self.module_list[i][0].anchors
        #Get the input dimensions
        inp_dim = int (self.net_info["height"])

        #Get the number of classes
        num_classes = int (module["classes"])

        #Transform
        x = x.data
        x = predict_transform(x, inp_dim, anchors, num_classes, CUDA)

        if not write: #if no collector has been initialised.

            detections = x
            write = 1

        else:
            detections = torch.cat((detections, x), 1)

    outputs[i] = x

```

```

        return detections

def load_weights(self, weightfile):
    #Open the weights file
    fp = open(weightfile, "rb")

    #The first 5 values are header information
    # 1. Major version number
    # 2. Minor Version Number
    # 3. Subversion number
    # 4,5. Images seen by the network (during training)
    header = np.fromfile(fp, dtype = np.int32, count = 5)
    self.header = torch.from_numpy(header)
    self.seen = self.header[3]

    weights = np.fromfile(fp, dtype = np.float32)

    ptr = 0
    for i in range(len(self.module_list)):
        module_type = self.blocks[i + 1]["type"]

        #If module_type is convolutional load weights
        #Otherwise ignore.

        if module_type == "convolutional":
            model = self.module_list[i]
            try:
                batch_normalize = int(self.blocks[i+1]["batch_n
ormalize"])
            except:
                batch_normalize = 0

            conv = model[0]

            if (batch_normalize):
                bn = model[1]

                #Get the number of weights of Batch Norm Layer
                num_bn_biases = bn.bias.numel()

                #Load the weights
                bn_biases = torch.from_numpy(weights[ptr:ptr +
num_bn_biases])
                ptr += num_bn_biases

                bn_weights = torch.from_numpy(weights[ptr: ptr
+ num_bn_biases])
                ptr += num_bn_biases

                bn_running_mean = torch.from_numpy(weights[ptr:
ptr + num_bn_biases])
                ptr += num_bn_biases

```

```

        bn_running_var = torch.from_numpy(weights[ptr:
ptr + num_bn_biases])
        ptr += num_bn_biases

        #Cast the loaded weights into dims of model wei
ghts.

        bn_biases = bn_biases.view_as(bn.bias.data)
        bn_weights = bn_weights.view_as(bn.weight.data)
        bn_running_mean = bn_running_mean.view_as(bn.ru
nning_mean)
        bn_running_var = bn_running_var.view_as(bn.runn
ing_var)

        #Copy the data to model
        bn.bias.data.copy_(bn_biases)
        bn.weight.data.copy_(bn_weights)
        bn.running_mean.copy_(bn_running_mean)
        bn.running_var.copy_(bn_running_var)

    else:
        #Number of biases
        num_biases = conv.bias.numel()

        #Load the weights
        conv_biases = torch.from_numpy(weights[ptr: ptr
+ num_biases])
        ptr = ptr + num_biases

        #reshape the loaded weights according to the di
ms of the model weights
        conv_biases = conv_biases.view_as(conv.bias.dat
a)

        #Finally copy the data
        conv.bias.data.copy_(conv_biases)

        #Let us load the weights for the Convolutional laye
rs

        num_weights = conv.weight.numel()

        #Do the same as above for weights
        conv_weights = torch.from_numpy(weights[ptr:ptr+num
_weights])
        ptr = ptr + num_weights

        conv_weights = conv_weights.view_as(conv.weight.dat
a)
        conv.weight.data.copy_(conv_weights)

```

```

In [9]: def predict_transform(prediction, inp_dim, anchors, num_classes, CUDA = True):

    batch_size = prediction.size(0)
    stride = inp_dim // prediction.size(2)
    grid_size = inp_dim // stride
    bbox_attrs = 5 + num_classes
    num_anchors = len(anchors)

    prediction = prediction.view(batch_size, bbox_attrs*num_anchors, grid_size*grid_size)
    prediction = prediction.transpose(1,2).contiguous()
    prediction = prediction.view(batch_size, grid_size*grid_size*num_anchors, bbox_attrs)
    anchors = [(a[0]/stride, a[1]/stride) for a in anchors]

    #Sigmoid the centre_X, centre_Y. and object confidence
    prediction[:, :, 0] = torch.sigmoid(prediction[:, :, 0])
    prediction[:, :, 1] = torch.sigmoid(prediction[:, :, 1])
    prediction[:, :, 4] = torch.sigmoid(prediction[:, :, 4])

    #Add the center offsets
    grid = np.arange(grid_size)
    a,b = np.meshgrid(grid, grid)

    x_offset = torch.FloatTensor(a).view(-1,1)
    y_offset = torch.FloatTensor(b).view(-1,1)

    if CUDA:
        x_offset = x_offset.cuda()
        y_offset = y_offset.cuda()

    x_y_offset = torch.cat((x_offset, y_offset), 1).repeat(1,num_anchors).view(-1,2).unsqueeze(0)

    prediction[:, :, :2] += x_y_offset

    #log space transform height and the width
    anchors = torch.FloatTensor(anchors)

    if CUDA:
        anchors = anchors.cuda()

    anchors = anchors.repeat(grid_size*grid_size, 1).unsqueeze(0)
    prediction[:, :, 2:4] = torch.exp(prediction[:, :, 2:4])*anchors

    prediction[:, :, 5: 5 + num_classes] = torch.sigmoid((prediction[:, :, 5 : 5 + num_classes]))

    prediction[:, :, :4] *= stride

    return prediction

```

```
In [10]: # wget https://github.com/ayoozhkathuria/pytorch-yolo-v3/raw/master
         /dog-cycle-car.png
```

```
In [11]: def get_test_input():
         img = cv2.imread("img/dog-cycle-car.png")
         img = cv2.resize(img, (416, 416))           #Resize to the input
         dimension
         img_ = img[:,:,:-1].transpose((2,0,1))      #BGR -> RGB | H X W C
         -> C X H X W
         img_ = img_[np.newaxis,:,:,:]/255.0        #Add a channel at 0 (
         for batch) | Normalise
         img_ = torch.from_numpy(img_).float()       #Convert to float
         img_ = Variable(img_)                       #Convert to Variable
         return img_
```

```
In [12]: model = Darknet("cfg/yolov3.cfg")
         inp = get_test_input()
         print(inp.shape)
         pred = model(inp, torch.cuda.is_available())
         print (pred)
```

```
torch.Size([1, 3, 416, 416])
tensor([[[[1.6675e+01, 1.4159e+01, 6.8440e+01, ..., 4.1797e-01,
          5.1057e-01, 5.2733e-01],
         [1.3412e+01, 1.5306e+01, 1.4919e+02, ..., 5.8069e-01,
          4.8104e-01, 5.1910e-01],
         [1.8618e+01, 1.9154e+01, 3.1965e+02, ..., 4.8069e-01,
          5.4098e-01, 5.5036e-01],
         ...,
         [4.1208e+02, 4.1248e+02, 1.1816e+01, ..., 5.2179e-01,
          5.2875e-01, 4.2780e-01],
         [4.1234e+02, 4.1198e+02, 1.7884e+01, ..., 5.7071e-01,
          4.9250e-01, 5.0524e-01],
         [4.1318e+02, 4.1198e+02, 5.2446e+01, ..., 3.9043e-01,
          4.6008e-01, 5.2686e-01]]]])
```

/Users/chanho/miniconda3/envs/yolo/lib/python3.8/site-packages/torch/nn/functional.py:2503: UserWarning: Default upsampling behavior when mode=bilinear is changed to align\_corners=False since 0.4.0. Please specify align\_corners=True if the old behavior is desired. See the documentation of nn.Upsample for details.

warnings.warn("Default upsampling behavior when mode={} is changed "

```
In [13]: # wget https://pjreddie.com/media/files/yolov3.weights
```

```
In [14]: model = Darknet("cfg/yolov3.cfg")
         model.load_weights("yolov3.weights")
```

```
In [15]: def bbox_iou(box1, box2):
    """
    Returns the IoU of two bounding boxes

    """
    #Get the coordinates of bounding boxes
    b1_x1, b1_y1, b1_x2, b1_y2 = box1[:,0], box1[:,1], box1[:,2], box1[:,3]
    b2_x1, b2_y1, b2_x2, b2_y2 = box2[:,0], box2[:,1], box2[:,2], box2[:,3]

    #get the coordinates of the intersection rectangle
    inter_rect_x1 = torch.max(b1_x1, b2_x1)
    inter_rect_y1 = torch.max(b1_y1, b2_y1)
    inter_rect_x2 = torch.min(b1_x2, b2_x2)
    inter_rect_y2 = torch.min(b1_y2, b2_y2)

    #Intersection area
    inter_area = torch.clamp(inter_rect_x2 - inter_rect_x1 + 1, min=0) * torch.clamp(inter_rect_y2 - inter_rect_y1 + 1, min=0)

    #Union Area
    b1_area = (b1_x2 - b1_x1 + 1)*(b1_y2 - b1_y1 + 1)
    b2_area = (b2_x2 - b2_x1 + 1)*(b2_y2 - b2_y1 + 1)

    iou = inter_area / (b1_area + b2_area - inter_area)

    return iou
```

## Part 4 : Objectness score thresholding and Non-maximum suppression (<https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-4/>)

```
In [16]: def write_results(prediction, confidence, num_classes, nms_conf = 0.4):
    conf_mask = (prediction[:, :, 4] > confidence).float().unsqueeze(2)
    prediction = prediction*conf_mask

    box_corner = prediction.new(prediction.shape)
    box_corner[:, :, 0] = (prediction[:, :, 0] - prediction[:, :, 2])/2
    box_corner[:, :, 1] = (prediction[:, :, 1] - prediction[:, :, 3])/2
    box_corner[:, :, 2] = (prediction[:, :, 0] + prediction[:, :, 2])/2
    box_corner[:, :, 3] = (prediction[:, :, 1] + prediction[:, :, 3])/2
    prediction[:, :, 4] = box_corner[:, :, 4]

    batch_size = prediction.size(0)
```

```

write = False

for ind in range(batch_size):
    image_pred = prediction[ind]           #image Tensor
    #confidence thresholding
    #NMS

    max_conf, max_conf_score = torch.max(image_pred[:,5:5+ num_
classes], 1)
    max_conf = max_conf.float().unsqueeze(1)
    max_conf_score = max_conf_score.float().unsqueeze(1)
    seq = (image_pred[:,5], max_conf, max_conf_score)
    image_pred = torch.cat(seq, 1)

    non_zero_ind = (torch.nonzero(image_pred[:,4]))
    try:
        image_pred_ = image_pred[non_zero_ind.squeeze(),:].view
(-1,7)
    except:
        continue

    if image_pred_.shape[0] == 0:
        continue
#

    #Get the various classes detected in the image
    img_classes = unique(image_pred_[:-1]) # -1 index holds t
he class index

    for cls in img_classes:
        #perform NMS

        #get the detections with one particular class
        cls_mask = image_pred_*(image_pred_[:-1] == cls).float
().unsqueeze(1)
        class_mask_ind = torch.nonzero(cls_mask[:,2]).squeeze(
)
        image_pred_class = image_pred_[class_mask_ind].view(-1,
7)

        #sort the detections such that the entry with the maxim
um objectness
        #confidence is at the top
        conf_sort_index = torch.sort(image_pred_class[:,4], des
cending = True )[1]
        image_pred_class = image_pred_class[conf_sort_index]
        idx = image_pred_class.size(0) #Number of detections

```

```

        for i in range(idx):
            #Get the IOUs of all boxes that come after the one
            we are looking at
            #in the loop
            try:
                ious = bbox_iou(image_pred_class[i].unsqueeze(0
            ), image_pred_class[i+1:])
            except ValueError:
                break

            except IndexError:
                break

            #Zero out all the detections that have IoU > treshh
old
            iou_mask = (ious < nms_conf).float().unsqueeze(1)
            image_pred_class[i+1:] *= iou_mask

            #Remove the non-zero entries
            non_zero_ind = torch.nonzero(image_pred_class[:,4])
.squeeze()
            image_pred_class = image_pred_class[non_zero_ind].v
iew(-1,7)

            batch_ind = image_pred_class.new(image_pred_class.size(
0), 1).fill_(ind) #Repeat the batch_id for as many detections
of the class cls in the image
            seq = batch_ind, image_pred_class

            if not write:
                output = torch.cat(seq,1)
                write = True
            else:
                out = torch.cat(seq,1)
                output = torch.cat((output,out))

        try:
            return output
        except:
            return 0

```

## [Part 5 : Designing the input and the output pipelines](https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-5/) [\(https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-5/\)](https://blog.paperspace.com/how-to-implement-a-yolo-v3-object-detector-from-scratch-in-pytorch-part-5/)



```
In [17]: img_dir = 'img'
det_dir = 'det'
batch_size = 1
confidence = 0.5
nms_thesh = 0.4
cfgfile = 'cfg/yolov3.cfg'
weightsfile = 'yolov3.weights'
start = 0
reso = 416
CUDA = torch.cuda.is_available()
```

```
In [18]: # mkdir data
# cd data
# wget https://raw.githubusercontent.com/ayoochkathuria/YOLO_v3_tutorial_from_scratch/master/data/coco.names
```

```
In [19]: def load_classes(namesfile):
fp = open(namesfile, "r")
names = fp.read().split("\n")[:-1]
return names
```

```
In [20]: num_classes = 80      #For COCO
classes = load_classes("data/coco.names")
```

```
In [21]: #Set up the neural network
print("Loading network.....")
model = Darknet(cfgfile)
model.load_weights(weightsfile)
print("Network successfully loaded")

model.net_info["height"] = reso
inp_dim = int(model.net_info["height"])
assert inp_dim % 32 == 0
assert inp_dim > 32

#If there's a GPU available, put the model on GPU
if CUDA:
    model.cuda()

#Set the model in evaluation mode
model.eval()
```

```
Loading network.....
Network successfully loaded
```

```
Out[21]: Darknet(
  (module_list): ModuleList(
    (0): Sequential(
      (conv_0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_0): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True)
```

```

    ine=True, track_running_stats=True)
        (leaky_0): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (1): Sequential(
      (conv_1): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
      (batch_norm_1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_1): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (2): Sequential(
      (conv_2): Conv2d(64, 32, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_2): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (3): Sequential(
      (conv_3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_3): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (4): Sequential(
      (shortcut_4): EmptyLayer()
    )
    (5): Sequential(
      (conv_5): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
      (batch_norm_5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_5): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (6): Sequential(
      (conv_6): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1),
bias=False)
      (batch_norm_6): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_6): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (7): Sequential(
      (conv_7): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (batch_norm_7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_7): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (8): Sequential(
      (shortcut_8): EmptyLayer()
    )
    (9): Sequential(
      (conv_9): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1),

```

```

bias=False)
    (batch_norm_9): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_9): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (10): Sequential(
    (conv_10): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (batch_norm_10): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_10): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (11): Sequential(
    (shortcut_11): EmptyLayer()
  )
  (12): Sequential(
    (conv_12): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (batch_norm_12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_12): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (13): Sequential(
    (conv_13): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batch_norm_13): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_13): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (14): Sequential(
    (conv_14): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (batch_norm_14): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_14): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (15): Sequential(
    (shortcut_15): EmptyLayer()
  )
  (16): Sequential(
    (conv_16): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batch_norm_16): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_16): LeakyReLU(negative_slope=0.1, inplace=True)
  )
  (17): Sequential(
    (conv_17): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (batch_norm_17): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (leaky_17): LeakyReLU(negative_slope=0.1, inplace=True)
  )

```

```
(18): Sequential(
  (shortcut_18): EmptyLayer()
)
(19): Sequential(
  (conv_19): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_19): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_19): LeakyReLU(negative_slope=0.1, inplace=True)
)
(20): Sequential(
  (conv_20): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_20): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_20): LeakyReLU(negative_slope=0.1, inplace=True)
)
(21): Sequential(
  (shortcut_21): EmptyLayer()
)
(22): Sequential(
  (conv_22): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_22): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_22): LeakyReLU(negative_slope=0.1, inplace=True)
)
(23): Sequential(
  (conv_23): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_23): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_23): LeakyReLU(negative_slope=0.1, inplace=True)
)
(24): Sequential(
  (shortcut_24): EmptyLayer()
)
(25): Sequential(
  (conv_25): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_25): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_25): LeakyReLU(negative_slope=0.1, inplace=True)
)
(26): Sequential(
  (conv_26): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_26): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_26): LeakyReLU(negative_slope=0.1, inplace=True)
)
(27): Sequential(
  (shortcut_27): EmptyLayer()
)
```

```

    )
    (28): Sequential(
      (conv_28): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_28): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_28): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (29): Sequential(
      (conv_29): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_29): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_29): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (30): Sequential(
      (shortcut_30): EmptyLayer()
    )
    (31): Sequential(
      (conv_31): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_31): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_31): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (32): Sequential(
      (conv_32): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_32): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_32): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (33): Sequential(
      (shortcut_33): EmptyLayer()
    )
    (34): Sequential(
      (conv_34): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_34): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_34): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (35): Sequential(
      (conv_35): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_35): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_35): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (36): Sequential(
      (shortcut_36): EmptyLayer()
    )
    (37): Sequential(

```

```
(conv_37): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
(batch_norm_37): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(leaky_37): LeakyReLU(negative_slope=0.1, inplace=True)
)
(38): Sequential(
  (conv_38): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_38): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_38): LeakyReLU(negative_slope=0.1, inplace=True)
)
(39): Sequential(
  (conv_39): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_39): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_39): LeakyReLU(negative_slope=0.1, inplace=True)
)
(40): Sequential(
  (shortcut_40): EmptyLayer()
)
(41): Sequential(
  (conv_41): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_41): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_41): LeakyReLU(negative_slope=0.1, inplace=True)
)
(42): Sequential(
  (conv_42): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_42): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_42): LeakyReLU(negative_slope=0.1, inplace=True)
)
(43): Sequential(
  (shortcut_43): EmptyLayer()
)
(44): Sequential(
  (conv_44): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_44): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_44): LeakyReLU(negative_slope=0.1, inplace=True)
)
(45): Sequential(
  (conv_45): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_45): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_45): LeakyReLU(negative_slope=0.1, inplace=True)
```

```
)
(46): Sequential(
  (shortcut_46): EmptyLayer()
)
(47): Sequential(
  (conv_47): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_47): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_47): LeakyReLU(negative_slope=0.1, inplace=True)
)
(48): Sequential(
  (conv_48): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_48): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_48): LeakyReLU(negative_slope=0.1, inplace=True)
)
(49): Sequential(
  (shortcut_49): EmptyLayer()
)
(50): Sequential(
  (conv_50): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_50): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_50): LeakyReLU(negative_slope=0.1, inplace=True)
)
(51): Sequential(
  (conv_51): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_51): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_51): LeakyReLU(negative_slope=0.1, inplace=True)
)
(52): Sequential(
  (shortcut_52): EmptyLayer()
)
(53): Sequential(
  (conv_53): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_53): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_53): LeakyReLU(negative_slope=0.1, inplace=True)
)
(54): Sequential(
  (conv_54): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_54): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_54): LeakyReLU(negative_slope=0.1, inplace=True)
)
(55): Sequential(
```

```
        (shortcut_55): EmptyLayer()
    )
    (56): Sequential(
      (conv_56): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_56): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_56): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (57): Sequential(
      (conv_57): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_57): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_57): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (58): Sequential(
      (shortcut_58): EmptyLayer()
    )
    (59): Sequential(
      (conv_59): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_59): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_59): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (60): Sequential(
      (conv_60): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_60): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_60): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (61): Sequential(
      (shortcut_61): EmptyLayer()
    )
    (62): Sequential(
      (conv_62): Conv2d(512, 1024, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (batch_norm_62): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_62): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (63): Sequential(
      (conv_63): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_63): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_63): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (64): Sequential(
      (conv_64): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
```



```

        (batch_norm_64): BatchNorm2d(1024, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
        (leaky_64): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (65): Sequential(
      (shortcut_65): EmptyLayer()
    )
    (66): Sequential(
      (conv_66): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1,
1), bias=False)
      (batch_norm_66): BatchNorm2d(512, eps=1e-05, momentum=0.1, a
ffine=True, track_running_stats=True)
      (leaky_66): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (67): Sequential(
      (conv_67): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1,
1), padding=(1, 1), bias=False)
      (batch_norm_67): BatchNorm2d(1024, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (leaky_67): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (68): Sequential(
      (shortcut_68): EmptyLayer()
    )
    (69): Sequential(
      (conv_69): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1,
1), bias=False)
      (batch_norm_69): BatchNorm2d(512, eps=1e-05, momentum=0.1, a
ffine=True, track_running_stats=True)
      (leaky_69): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (70): Sequential(
      (conv_70): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1,
1), padding=(1, 1), bias=False)
      (batch_norm_70): BatchNorm2d(1024, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (leaky_70): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (71): Sequential(
      (shortcut_71): EmptyLayer()
    )
    (72): Sequential(
      (conv_72): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1,
1), bias=False)
      (batch_norm_72): BatchNorm2d(512, eps=1e-05, momentum=0.1, a
ffine=True, track_running_stats=True)
      (leaky_72): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (73): Sequential(
      (conv_73): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1,
1), padding=(1, 1), bias=False)
      (batch_norm_73): BatchNorm2d(1024, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)

```

```

        (leaky_73): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (74): Sequential(
      (shortcut_74): EmptyLayer()
    )
    (75): Sequential(
      (conv_75): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_75): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_75): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (76): Sequential(
      (conv_76): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_76): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_76): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (77): Sequential(
      (conv_77): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_77): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_77): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (78): Sequential(
      (conv_78): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_78): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_78): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (79): Sequential(
      (conv_79): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_79): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_79): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (80): Sequential(
      (conv_80): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_80): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_80): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (81): Sequential(
      (conv_81): Conv2d(1024, 255, kernel_size=(1, 1), stride=(1, 1))
    )
    (82): Sequential(
      (Detection_82): DetectionLayer()
    )

```

```
)
(83): Sequential(
  (route_83): EmptyLayer()
)
(84): Sequential(
  (conv_84): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_84): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_84): LeakyReLU(negative_slope=0.1, inplace=True)
)
(85): Sequential(
  (upsample_85): Upsample(scale_factor=2.0, mode=bilinear)
)
(86): Sequential(
  (route_86): EmptyLayer()
)
(87): Sequential(
  (conv_87): Conv2d(768, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_87): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_87): LeakyReLU(negative_slope=0.1, inplace=True)
)
(88): Sequential(
  (conv_88): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_88): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_88): LeakyReLU(negative_slope=0.1, inplace=True)
)
(89): Sequential(
  (conv_89): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_89): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_89): LeakyReLU(negative_slope=0.1, inplace=True)
)
(90): Sequential(
  (conv_90): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (batch_norm_90): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_90): LeakyReLU(negative_slope=0.1, inplace=True)
)
(91): Sequential(
  (conv_91): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batch_norm_91): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (leaky_91): LeakyReLU(negative_slope=0.1, inplace=True)
)
(92): Sequential(
```

```

        (conv_92): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (batch_norm_92): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (leaky_92): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (93): Sequential(
      (conv_93): Conv2d(512, 255, kernel_size=(1, 1), stride=(1, 1))
    )
    (94): Sequential(
      (Detection_94): DetectionLayer()
    )
    (95): Sequential(
      (route_95): EmptyLayer()
    )
    (96): Sequential(
      (conv_96): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_96): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_96): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (97): Sequential(
      (upsample_97): Upsample(scale_factor=2.0, mode=bilinear)
    )
    (98): Sequential(
      (route_98): EmptyLayer()
    )
    (99): Sequential(
      (conv_99): Conv2d(384, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_99): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_99): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (100): Sequential(
      (conv_100): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (batch_norm_100): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_100): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (101): Sequential(
      (conv_101): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (batch_norm_101): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (leaky_101): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (102): Sequential(
      (conv_102): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

```

```

        (batch_norm_102): BatchNorm2d(256, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
        (leaky_102): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (103): Sequential(
      (conv_103): Conv2d(256, 128, kernel_size=(1, 1), stride=(1,
1), bias=False)
      (batch_norm_103): BatchNorm2d(128, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (leaky_103): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (104): Sequential(
      (conv_104): Conv2d(128, 256, kernel_size=(3, 3), stride=(1,
1), padding=(1, 1), bias=False)
      (batch_norm_104): BatchNorm2d(256, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (leaky_104): LeakyReLU(negative_slope=0.1, inplace=True)
    )
    (105): Sequential(
      (conv_105): Conv2d(256, 255, kernel_size=(1, 1), stride=(1,
1))
    )
    (106): Sequential(
      (Detection_106): DetectionLayer()
    )
  )
)

```

```

In [22]: read_dir = time.time()
         #Detection phase
         try:
             imlist = [osp.join(osp.realpath('.'), img_dir, img) for img in
os.listdir(img_dir) if img.find("png") > -1 ]
         except NotADirectoryError:
             imlist = []
             imlist.append(osp.join(osp.realpath('.'), img_dir))
         except FileNotFoundError:
             print ("No file or directory with the name {}".format(img_dir))
             exit()

```

```

In [23]: txt = '/Users/chanho/Documents/GitHub/keras-YOLOv3/img/frame0700.pn
g'
         print(txt.find("png") > -1)

```

True

```
In [24]: print(imlist)

['/Users/chanho/Documents/GitLab/niceface/models/img/frame0900.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/zebra.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/frame1700.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/frame1900.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/frame1309.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/frame0600.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/frame1291.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/frame1441.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/dog-cycle-car.png',
 '/Users/chanho/Documents/GitLab/niceface/models/img/frame0700.png']
```

```
In [25]: if not os.path.exists(det_dir):
         os.makedirs(det_dir)
```

```
In [26]: load_batch = time.time()
         loaded_ims = [cv2.imread(x) for x in imlist]
```

```
In [27]: def letterbox_image(img, inp_dim):
         '''resize image with unchanged aspect ratio using padding'''
         img_w, img_h = img.shape[1], img.shape[0]
         w, h = inp_dim
         new_w = int(img_w * min(w/img_w, h/img_h))
         new_h = int(img_h * min(w/img_w, h/img_h))
         resized_image = cv2.resize(img, (new_w,new_h), interpolation =
cv2.INTER_CUBIC)

         canvas = np.full((inp_dim[1], inp_dim[0], 3), 128)

         canvas[(h-new_h)//2:(h-new_h)//2 + new_h,(w-new_w)//2:(w-new_w)
//2 + new_w, :] = resized_image

         return canvas
```

```
In [28]: def prep_image(img, inp_dim):
         """
         Prepare image for inputting to the neural network.

         Returns a Variable
         """
         img = (letterbox_image(img, (inp_dim, inp_dim)))
         img = img[:, :, :-1].transpose((2,0,1)).copy()
         img = torch.from_numpy(img).float().div(255.0).unsqueeze(0)
         return img
```

```
In [29]: #PyTorch Variables for images
im_batches = list(map(prepare_image, loaded_imgs, [inp_dim for x in range(len(imlist))]))

#List containing dimensions of original images
im_dim_list = [(x.shape[1], x.shape[0]) for x in loaded_imgs]
im_dim_list = torch.FloatTensor(im_dim_list).repeat(1,2)

if CUDA:
    im_dim_list = im_dim_list.cuda()
```

```
In [30]: leftover = 0
if (len(im_dim_list) % batch_size):
    leftover = 1

if batch_size != 1:
    num_batches = len(imlist) // batch_size + leftover
    im_batches = [torch.cat((im_batches[i*batch_size : min((i + 1)*
batch_size,
                                len(im_batches))])) for i in range(num_batches)]
```

```
In [31]: def unique(tensor):
    tensor_np = tensor.cpu().numpy()
    unique_np = np.unique(tensor_np)
    unique_tensor = torch.from_numpy(unique_np)

    tensor_res = tensor.new(unique_tensor.shape)
    tensor_res.copy_(unique_tensor)
    return tensor_res
```

```

In [32]: write = 0
start_det_loop = time.time()
for i, batch in enumerate(im_batches):
    #load the image
    start = time.time()
    if CUDA:
        batch = batch.cuda()

    prediction = model(Variable(batch, volatile = True), CUDA)

    prediction = write_results(prediction, confidence, num_classes,
nms_conf = nms_thesh)

    end = time.time()

    if type(prediction) == int:

        for im_num, image in enumerate(imlist[i*batch_size: min((i
+ 1)*batch_size, len(imlist))]):
            im_id = i*batch_size + im_num
            print("{0:20s} predicted in {1:6.3f} seconds".format(im
age.split("/")[-1], (end - start)/batch_size))
            print("{0:20s} {1:s}".format("Objects Detected:", ""))
            print("-----")
        continue

    prediction[:,0] += i*batch_size    #transform the attribute from
index in batch to index in imlist

    if not write:                    #If we have't initialised ou
tput
        output = prediction
        write = 1
    else:
        output = torch.cat((output,prediction))

    for im_num, image in enumerate(imlist[i*batch_size: min((i + 1
)*batch_size, len(imlist))]):
        im_id = i*batch_size + im_num
        objs = [classes[int(x[-1])] for x in output if int(x[0]) ==
im_id]
        print("{0:20s} predicted in {1:6.3f} seconds".format(image.
split("/")[-1], (end - start)/batch_size))
        print("{0:20s} {1:s}".format("Objects Detected:", " ".join(
objs)))
        print("-----")
    print("-----")

    if CUDA:
        torch.cuda.synchronize()

```



```

<ipython-input-32-e5724478fd54>:9: UserWarning: volatile was removed and now has no effect. Use `with torch.no_grad():` instead.
    prediction = model(Variable(batch, volatile = True), CUDA)

frame0900.png          predicted in  0.863 seconds
Objects Detected:      person person person person person person person
son person
-----
zebra.png              predicted in  0.801 seconds
Objects Detected:      zebra zebra zebra
-----
frame1700.png          predicted in  0.782 seconds
Objects Detected:      person person person person person
-----
frame1900.png          predicted in  0.774 seconds
Objects Detected:      person person person person person person person
son person person
-----
frame1309.png          predicted in  0.771 seconds
Objects Detected:      person person person person person person
-----
frame0600.png          predicted in  0.795 seconds
Objects Detected:      person person person person person person person
son train backpack
-----
frame1291.png          predicted in  0.808 seconds
Objects Detected:      person person person
-----
frame1441.png          predicted in  0.813 seconds
Objects Detected:      person person person person person
-----
dog-cycle-car.png      predicted in  0.830 seconds
Objects Detected:      bicycle truck dog
-----
frame0700.png          predicted in  0.780 seconds
Objects Detected:      person person person person
-----

```

```

In [33]: try:
          output
        except NameError:
          print ("No detections were made")
          exit()

```

```
In [34]: im_dim_list = torch.index_select(im_dim_list, 0, output[:,0].long()
        )

        scaling_factor = torch.min(inp_dim/im_dim_list,1)[0].view(-1,1)

        output[:,[1,3]] -= (inp_dim - scaling_factor*im_dim_list[:,0].view(
        -1,1))/2
        output[:,[2,4]] -= (inp_dim - scaling_factor*im_dim_list[:,1].view(
        -1,1))/2

        output[:,1:5] /= scaling_factor
```

```
In [35]: for i in range(output.shape[0]):
        output[i, [1,3]] = torch.clamp(output[i, [1,3]], 0.0, im_dim_li
        st[i,0])
        output[i, [2,4]] = torch.clamp(output[i, [2,4]], 0.0, im_dim_li
        st[i,1])
```

```
In [36]: class_load = time.time()
        colors = pickle.load(open("data/pallete", "rb"))
```

```
In [37]: draw = time.time()

        def write(x, results, color):
            c1 = tuple(x[1:3].int())
            c2 = tuple(x[3:5].int())
            img = results[int(x[0])]
            cls = int(x[-1])
            label = "{0}".format(classes[cls])
            cv2.rectangle(img, c1, c2,color, 1)
            t_size = cv2.getTextSize(label, cv2.FONT_HERSHEY_PLAIN, 1 , 1)[
            0]
            c2 = c1[0] + t_size[0] + 3, c1[1] + t_size[1] + 4
            cv2.rectangle(img, c1, c2,color, -1)
            cv2.putText(img, label, (c1[0], c1[1] + t_size[1] + 4), cv2.FON
            T_HERSHEY_PLAIN, 1, [225,255,255], 1);
            return img
```

```
In [38]: list(map(lambda x: write(x, loaded_ims, colors[0]), output))
        det_names = pd.Series(imlist).apply(lambda x: "{}/det_{}".format(de
        t_dir,x.split("/")[-1]))
        list(map(cv2.imwrite, det_names, loaded_ims))
        end = time.time()
```

Here is a sample image of results.

